

## CS4023 Week04 Lab Exercise

**Lab Objective:** We will use this lab to log in to our Linux Systems and look at some simple programs that perform a few elementary system calls. By the end of the lab, we will have developed the machinery to more or less mimic UNIX's cp copy command.

Here's a quick summary of the tasks:

1. Create a series of hierarchical directories to manage our work throughout the remainder of the semester
2. Copy the C source code readFile.c into your directory and compile it
3. Watch it in action, tracing “system events” using the strace and ptrace utilities
4. Learn some of the basic tricks of a good programmer
5. Extend the given program so that it reads a file and writes the contents somewhere else
6. Consider what is wrong with the program we have just worked on.

### **In Detail**

Note that there are a few alternatives given below for this first step so please read the Entire set of instructions before acting. Over the semester we will have labs, programming assignments, etc. and it is a good idea to keep these in an organised way. My suggestion is that you create a subdirectory on your home directory called cs4023 that will be the top-level point for all module-related material. The command to create a directory in Linux is mkdir, so you could do

```
mkdir ~/cs4023
```

which makes a subdirectory located in your home directory. Next you should make a subdirectory called labs in this for each week's work. This can be done with

```
mkdir ~/cs4023/labs
```

Finally, for this week's lab, Week04, you should create its own subdirectory with

```
mkdir ~/cs4023/labs/week04
```

An alternative to making each level of the hierarchy at a time is to tell mkdir to make the “parent” subdirectory if it doesn't exist. So the following command can take the place of all of the previous ones.

```
mkdir -p ~/cs4023/labs/week04
```

The -p is for “make parents if they don't already exist.” Note that it is very similar to the command before it, but you had to do a lot of extra work in order to achieve that. Yet another alternative is to bring up a GUI-based file manager and use the “Create ...” menu option there.

Using the C code provided in the class directory, (Sulis->Week4), copy this file into your directory ~/cs4023/labs/week04. This could be achieved using wget. This is an unguided challenge in this lab, figure out an efficient way of copying the code from Sulis to Linux.

Here's something that gets people every time. The tilde character, ~, has a special use at the CLI (command-line interface). Several different shells { including bash, which we are using { treat the tilde character as denoting a user's home directory. So, no matter where you are in your own filesystem, you can use ~ as a shorthand for your home directory, which, for user d9994023 will be /home/d9994023. But you can also use ~ to refer to somebody else's home directory. In the previous command ~cs4023 is a shorthand for the home directory of user cs4023. So, there is a huge difference between ~cs4023 and ~/cs4023: the former is the home directory of user cs4023 while the latter is a sub-directory called cs4023 of your home directory. In summary: watch out for not confusing the two arguments to the cp command!

Once this is done you can examine the source code in a text editor. The one I use is emacs, which may or not be installed on your lab machine. If it is, you can fire it up with

```
emacs readFile.c &
```

but you need to be in the correct directory for this. As an alternative editor many people in the past have preferred gedit, but this is subject to your machines. You can change working directory with the command

```
cd ~/cs4023/labs/week04
```

The ampersand at the end is important for it runs the program in the background, allowing you to give further commands at the shell prompt. Please spend a while looking at the code; it shouldn't all make sense to you but between looking at the source and running it (later), you should get an idea of what it does. One thing to look at is the while loop. As you know from your Java days while executes the body of the while-loop as long as some condition is true. In this case the "condition" is decided by comparing the result of "c = getc(from)" to the special character EOF. This special character exists at the end of every UNIX file and is used, for example, by a disk-reading program to tell when to stop reading. As long as c is not this special character the while-loop will keep on running. The character is "put" to stdout, which is the screen in this case.

Now compile this program with the command

```
gcc -o readFile readFile.c
```

Actually, the command gcc does two jobs: it firstly compiles the file readFile.c into CPU-specific code (machine or object code) and then links together this object code with any necessary support modules of code to create an executable program.

Once you have compiled your program you should have an executable called readFile that you can run as follows:

```
readFile
```

If you have been following along with the previous labs, doing as was suggested, then the above should work. If it doesn't work it is likely that you have not modified your `$PATH` environment variable to instruct the shell to look in the current working directory first for commands to execute. This environment variable should be in either of the files `~/.bashrc` or `~/.bash_profile` in your home directory.

You can avoid always having to specify, `./`, the current working directory, by a quick edit of your `.bash` profile in your home directory. In this file you will find a line that begins with `"PATH="`. What follows is a list of directories, each separated from the next by a `:`. Whenever you issue a command, this list of directories or locations is where the shell looks to try and find an executable program. So, by putting just a simple `.` at the end of the line will now tell the shell interpreter to look in the current directory as a last resort for a command to execute.

- 1) Don't forget the `:` separator and,
- 2) this will work next time you log in; for the present session, issue the command `sh $HOME/.bash profile`.

When you do get the program to run, it should report back immediately with `file.txt: No such file or directory`. This is because we don't have such a file in this directory. In order to keep `readFile` happy please create a file called `file.txt` now. You can edit this from scratch with anything you like or you can do something like

```
head -5 readFile.c > file.txt
```

which will insert the first 5 lines of your C source file into the file `file.txt`. If parsimoniousness is your concern you could even do

```
touch file.txt
```

(Even though we mentioned it in class recently this one is worth doing a man on. Do it now. That is, at the shell prompt give the command `man touch`. Note that order is important, this is not the same as `touch man`.) :P

Once you have created something and called it `file.txt` you should be able to run `readFile` and have it exit successfully. Now we will cause the program to fail, and this will demonstrate the power of the `perror()` function call in the program. Change the protections on the file `file.txt` so that it is not readable. You can do this with

```
chmod 0333 file.txt
```

or

```
chmod ugo-r file.txt
```

Now run `readFile` again. Note the different error message. When attempting to open the file `fopen()` returned `NULL` and we then had `perror()` decode the reason why it failed. Please, please get in the habit of testing the result of every system call you make - `fopen()` in this

case - and reporting the error with `perror()` in the case of failure. Note also the `fclose()` system call that closes the file opened for reading. Just like your mother told you to always close gates that you opened, so too should you `fclose` files `fopened` by you.

Don't forget to change back to something sensible the protections on the file `file.txt` now.

A very useful tool to help in debugging programs that make system calls or otherwise interact with the kernel is the `strace` program. This program and its (simpler and less verbose) counterpart, `ltrace`, allow you to see a report of system calls and library calls - `strace` is for system calls and `ltrace` is for library calls.

They do their job by you giving your command as normal and just prepending either one or the other. For example

`ltrace readFile`

reports the library calls made in running our program. It is the program `ltrace` that is executed first and this manages the running of our program `readFile`.

You will now write a new program called `copyFile` that will read the file `file.txt` and now write the contents to the file `copy.txt`. Firstly, make a copy of the `readFile.c` source code and save it in `copyFile.c`. Now edit this file so that the copying of `file.txt` gets done. This involves two main steps:

- opening - and closing - two files instead of one
- writing each byte read from `file.txt` to `copy.txt`

You will need to look at the syntax of the `fopen()` system call to find out how to open a file for writing. Do `man fopen` to learn this.

The second item will require you to modify the body of the while loop in the program so that the character is no longer put to standard output (`stdout`). You will need to put it to the destination file instead.

This is the first lab that you will be assessed/graded on. There will be 6 assessments in total, each worth 5%. To be assessed with no penalty charged you will need to have the program written, working properly, and handed in by 18.00, Thursday. Week05.

So, over the semester, when you are given a programming assignment, please make sure that the output of your program matches exactly with the output you were asked to generate.

You will get 1 Week to work on each Lab Assignment in this module

Closing comments:

The trouble with this `copyFile` program is its inflexibility: it is impossible to read from a file other than `file.txt` or to write it to anything other than `copy.txt`. What we would like to be able to do is to specify the name of the file to read as a part of the command. This is where command-line arguments come in to play and we will discuss this sometime in the lab slot next time.